

Scikit-learn

- Simple and efficient tools for data mining and data analysis
- Accessible to everybody, and reusable in various contexts
- Built on NumPy, SciPy, and matplotlib
- Open source, commercially usable




Classification

Identifying which category an object belongs to.

Applications: Spam detection, image recognition.

Algorithms: [SVM](#), [nearest neighbors](#), [random forest](#), and [more...](#)




Regression

Predicting a continuous-valued attribute associated with an object.

Applications: Drug response, Stock prices.

Algorithms: [SVR](#), [nearest neighbors](#), [random forest](#), and [more...](#)




Clustering

Automatic grouping of similar objects into sets.

Applications: Customer segmentation, Grouping experiment outcomes

Algorithms: [k-Means](#), [spectral clustering](#), [mean-shift](#), and [more...](#)




Dimensionality reduction

Reducing the number of random variables to consider.

Applications: Visualization, Increased efficiency

Algorithms: [k-Means](#), [feature selection](#), [non-negative matrix factorization](#), and [more...](#)




Model selection

Comparing, validating and choosing parameters and models.

Applications: Improved accuracy via parameter tuning

Algorithms: [grid search](#), [cross validation](#), [metrics](#), and [more...](#)




Preprocessing

Feature extraction and normalization.

Applications: Transforming input data such as text for use with machine learning algorithms.

Algorithms: [preprocessing](#), [feature extraction](#), and [more...](#)



Data in scikit-learn stored as a 2D array `[m_samples, n_features]`

Scikit-learn's interface

Linear Models

```
from sklearn.linear_model import LinearRegression  
model = LinearRegression()  
model.fit(X_train, Y_train)  
model.predict(X_new)
```

Ordinary least squares
Linear Regression

Stochastic gradient descent

```
from sklearn.linear_model import SGDRegressor  
model = SGDRegressor(loss='squared_loss', penalty='l2',  
                      alpha=0.0001, learning_rate='constant', eta0=0.01)
```

$$J(\theta) = MSE(\theta) + \frac{\lambda}{2} \sum_{i=1}^n \theta_i^2 \quad (\text{Lc 1})$$

```
from sklearn.linear_model import LogisticRegression  
model = LogisticRegression(solver='lbfgs', max_iter=100, multi_class='auto')
```

solver{*'newton-cg'*, *'lbfgs'*, *'liblinear'*, *'sag'*, *'saga'*}, *default='lbfgs'*
Algorithm to use in the optimization problem

kNN

```
sklearn.neighbors.KNeighborsClassifier(n_neighbors=5, weights='uniform', p=2,  
                                       metric='minkowski')
```

```
sklearn.neighbors.KNeighborsRegressor()
```

Decision Tree

```
sklearn.tree.DecisionTreeClassifier(criterion='gini', max_depth=None,  
                                    min_samples_split=2, class_weight=None)
```

```
sklearn.tree.DecisionTreeRegressor()
```

Scikit-learn

Model Evaluation

```
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
```

```
y_true = [0, 0, 0, 1, 1, 2, 2, 2]
y_pred = [0, 0, 1, 1, 1, 2, 1, 2]
print(classification_report(y_true, y_pred))
```


	precision	recall	f1-score	support
0	1.00	0.67	0.80	3
1	0.50	1.00	0.67	2
2	1.00	0.67	0.80	3
accuracy			0.75	8
macro avg	0.83	0.78	0.76	8
weighted avg	0.88	0.75	0.77	8

Train Test split

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.30)
```


```
conf = confusion_matrix(y_true, y_pred)
print(conf)
[[ 2  1  0]
 [ 0  2  0]
 [ 0  1  2]]
```

```
import seaborn as sns
sns.heatmap(conf, annot=True)
```



Scikit-learn

Data Normalization



```
sklearn.preprocessing.StandardScaler()
```

Standardize features by removing the mean and scaling to unit variance

```
sklearn.preprocessing.MinMaxScaler(feature_range=0, 1)
```

Transform features by scaling each feature to a given range

Usage:

`fit(X)`

`transform(X)`

or

`fit_transform(X)`

SVM in sklearn lib

```
sklearn.svm.LinearSVC()  
sklearn.linear_model.SGDClassifier(loss='hinge')  
sklearn.svm.SVC(kernel='rbf')
```


Class	Time complexity	Scaling required	Kernel trick
LinearSVC	$O(m \times n)$	Yes	No
SGDClassifier	$O(m \times n)$	Yes	No
SVC	$O(m^2 \times n)$ to $O(m^3 \times n)$	Yes	Yes

SVMs are sensitive to the feature scaling!!!


SVM software packages:

- libsvm – most commonly used implementation of kernalized svm, sklearn uses wrapper over it
- liblinear – gradient descent based implementation of linear SVM


SVM Example




SVC(kernel='linear', C=1)




SVC(kernel='rbf', gamma=.01, C=1)




SVC(kernel='rbf', gamma=1, C=1)




SVC(kernel='rbf', gamma=10, C=1)




SVC(kernel='rbf', gamma=100, C=1)




SVM Example




SVC(kernel='rbf', gamma=.01, C=1)




SVC(kernel='rbf', gamma=.01, C=10)




SVC(kernel='rbf', gamma=.01, C=1000)



SVC(kernel='rbf', gamma=.01, C=10000)



SVC(kernel='rbf', gamma=.01, C=100000)



SVM in sklearn lib

Regression:

```
sklearn.svm.LinearSVR()  
Sklearn.svm.SVR()
```

GridSearch:

```
param_grid = {'C': [0.1,1, 10, 100, 1000], 'gamma': [1,0.1,0.01,0.001,0.0001], 'kernel': ['rbf']}  
from sklearn.model_selection import GridSearchCV  
grid = GridSearchCV(SVC(),param_grid,refit=True,verbose=3)  
grid.fit(X,y)  
print(grid.best_params_)
```

Ensemble Classifiers

`sklearn.ensembles.AdaBoostClassifier`


`sklearn.ensembles.RandomForestClassifier`

Feature importance with Random Forest:

- A single DT: important features are likely to appear closer to the root of tree, while unimportant features often appear closer to the leaves
- RF: feature importance as average depth at which it appears across all trees in the forest

`RandomForestClassifier().feature_importances_`

MNIST pixel importance with accordance to RF classifier



The MNIST database is a large [database](#) of handwritten digits (28x28 pixels, 60,000 training images and 10,000 testing images)

